

IIM-harness

python 3.12

PyTorch 2.5.1+cu121

provider modern-transformer


status documentation first

license MIT

contributions welcome

IIM-harness

the loop, the tools and the context engineering that turn a language model into an agent



Project Description

A language model is **stateless**. It reads tokens, it writes a few more, and it forgets everything. It cannot open a file, run a command, or remember what it did thirty seconds ago. Everything people call "an agent" lives in the program wrapped around that forward pass, and that program is the **harness**.

This repository documents how that harness is built, from the public post-mortems of the **March 2026 Claude Code source leak**, the clean-room rewrites published afterwards, and a direct reading of a running harness from the inside.

The target provider is my own [modern-transformer](#), so the same loop drives a local model on a 6 GB GPU or a hosted API without changing a line of the loop.

 **The documentation is the deliverable right now.** The Python package in "Repository structure" is the plan, not shipped code.



Features



The loop in fifteen lines, and why the quality is never in the loop itself



Tool design rules, few and wide, descriptions written as prompts, errors returned as results



Context engineering, the cache boundary, the append-only history, and disclosure on demand



Delegation, subagents, skills and hooks, plus when each one is the wrong answer



The full command intake, what the harness answers alone and what actually reaches the model



Provider-agnostic, the loop never knows if it talks to a local transformer or an API

Example Outputs

One real loop, printed the way a harness log reads. The model never sees the repository, only what the tools bring back.

```
turn 1  user      "why does the training loss explode around step 400"
turn 1  model      stop_reason = tool_use
        harness    └─ Grep(pattern="warmup", path="src/mt")      ┌─ dispatched
        harness    └─ Read("configs/base.yaml")                 └─ in parallel (340 ms)

turn 2  model      stop_reason = tool_use
        harness    └─ Read("src/mt/optim.py", offset=120, limit=60) (12 ms)

turn 3  model      stop_reason = tool_use
        harness    └─ Read("outputs/logs/base_run-07/run.log", offset=380)
        harness    ERROR: file not found                          (3 ms)

turn 4  model      stop_reason = tool_use
        harness    └─ Glob("outputs/logs/**/run.log")              (18 ms)

turn 5  model      stop_reason = end_turn
        model      "warmup ends at step 400 and the schedule hands over to
                   a cosine decay that restarts at the peak LR ..."
```

Turn 3 is the interesting one. The model guessed a path, it did not exist, and the harness returned the error as a normal result instead of crashing. The model read it, switched to `Glob`, and recovered on its own. Nobody wrote that recovery, it falls out of the loop.

Notes & Observations

 **Latency comes from the model, not the tools.** Four tool calls cost 373 ms together. Five forward passes cost several seconds. That is why independent calls are dispatched at once.

 **What the cache actually holds.** Not the text, the attention state: every token's key and value vectors, computed once and kept in the provider's VRAM, indexed by the exact token prefix that produced them. Each call still re-sends the whole conversation, the provider matches the longest identical prefix, reloads those vectors instead of recomputing them, and only the new tail goes through the model. **One changed byte and the match stops right there**, which is why a clock in the prefix is fatal. Entries expire on a short timer (five minutes by default, an hour on request), so it is a hot-path optimisation and never storage.

 **Five turns means five full API calls**, each one re-sending the whole conversation. Counting the tokens each turn puts in front of the model:

Turn	Conversation size	Billed without cache	Billed with cache
1	14 200	14 200	14 200
2	15 100	15 100	900
3	15 400	15 400	300
4	15 900	15 900	500
5	16 700	16 700	800
Total paid		77 300	16 700

Read the last row. Without a cache you pay 77 300 tokens for a conversation that only ever contained 16 700, because every turn re-bills the whole history. With a cache **the total cost of a session equals the final size of its window**, since each turn only bills what it added. Quadratic becomes linear, and that single row is why long agent sessions are affordable at all.

How it works

 **The harness assembles a context window.** A stable prefix (identity, rules, tool schemas) then the volatile tail (the conversation).

 **The model does one forward pass.** No memory of the previous turn, it re-reads the whole story every time.

 **The harness reads** `stop_reason` . That one field is the entire control flow.

 **Tools run**, and independent calls run at the same time.

 **Results are appended** as a new turn, shaped exactly like a user message.


 **The loop runs again**, until the model stops asking or a budget is hit.

The whole trick is that **the model decides and the harness executes**, and neither ever does the other's job.

Architecture Diagram

Agentic harness architecture

The model is one box. Everything around it is the harness, and that is where the behaviour comes from.



The centre column is the loop, small on purpose. The left column is everything competing for room in the window. The right column keeps a long session alive. The four mechanisms below cover them in order.

1 The loop

Every harness has this, and nobody is impressed by it.

```
messages = [user_turn]

while True:
    response = provider.complete(system=system_prompt, messages=messages, tools=tool_schemas)
    messages.append(response)

    if response.stop_reason != "tool_use":
        return response.text  # end_turn, done

    results = executor.run_all(response.tool_calls)  # independent calls in parallel
    messages.append(as_user_turn(results))  # results become the next turn
```

The quality difference is not here. It is in the two objects this loop manipulates, the tools it exposes and the window it builds.

- 🔧 **A failing tool returns a normal result containing the error.** No raise, no retry policy, no abort. The harness has zero recovery logic, it just refuses to hide failure from the one component able to reason about it.
- ⚡ **Parallelism is a prompt, not a scheduler.** The harness cannot know which calls are independent, the model can, so the rule lives in the prompt and the executor dispatches whatever arrived together.

2 The tools

A tool description is not documentation, it is instruction text the model re-reads every turn.

```
Bash = Tool(  
    description="""Executes a bash command and returns its output.  
  
    - IMPORTANT: Avoid using this tool to run find, grep, cat, head, tail or sed,  
    unless explicitly instructed. Use the dedicated tool instead.  
    - Interactive flags (-i) are not supported in this environment."""  
)
```

Those two lines are behaviour, not API reference. They sit in the cached prefix, so they are the cheapest place in the whole system to put a rule.

📦 **Few tools, wide.** One `Bash` beats twenty command wrappers, one `Edit` beats `insert_line` plus `delete_line` plus `replace_range`. Every schema costs tokens permanently.

🚫 **Fail loudly, never guess.** `Edit` demands the exact string and refuses a non-unique match. A silent wrong guess costs ten turns, a clean error costs one.

🔧 **MCP adds tools without touching the loop.** At startup the harness asks each declared server for its tool list and drops the schemas into the same registry. The model cannot tell an MCP tool from a native one, so adding a capability becomes a config entry instead of a code change.

3 The window

The window is finite, it is re-sent in full every turn, and every token competes with every other.

┌ stable prefix ───────────┐	
identity, rules, environment	cached, paid once
tool schemas, skill index, CLAUDE.md	
└─ cache boundary ─────────┘	← one timestamp above this line
turn 1 ... turn 47	and the cache misses every turn
tool results, injected reminders	paid every turn
└─ volatile tail ─────────┘	

📦 **The history is append-only, and that is a hard constraint.** A KV cache is only valid over a prefix, so changing a token at position `k` invalidates everything after it. Rewriting an old tool result costs more than the room it frees, so the harness never goes back and edits.

✂️ **It shapes output on the way in instead.** `Read` caps at 2 000 numbered lines and hands you `offset`, `Grep` stops at 250 matches, `WebFetch` runs the page through a small fast model so 50 000 tokens of HTML become 300 tokens of answer. Once inserted, that text is frozen forever.



Disclosure on demand, the same idea at three scales. Keep a catalogue in the window, pay for the content only when it is used:

Level	Always loaded	Loaded on demand
Skills	one description line each	the <code>SKILL.md</code> body, when invoked
Tool schemas	the tool name only	the full JSON schema, fetched by name
Memory	an index, one line per fact	the fact's own file, when relevant

Compaction is the one exception, and it is expensive: replacing the head of the conversation invalidates the whole cache, which is why it fires on a token budget and never continuously.

The pause that does not need to happen. A transformer *can* be paused and resumed, the KV cache is literally that state, and `mt.cache` implements it. A cloud GPU does not wait for your tool because VRAM is the scarce resource, not because the maths forbids it. Run the model locally and that trade-off flips, the cache can simply stay resident.

4 The delegation

A subagent is a second loop with its own window, and only its conclusion comes back.

```
---
name: ablation-runner
description: >
  Runs one modern-transformer ablation and reports the metric delta.
  Use when asked to compare two configs.
tools: Read, Grep, Glob, Bash
model: sonnet
---
```

You run a single ablation. Read the two configs, launch the training script, follow the log, and report only the run number and the metric delta.

Delegation is compression. Reading forty files costs 90 000 tokens inside the subagent and 400 in the parent. A task that emits as much as it consumes gains nothing, which is why `Explore` is the one that actually gets used and "write this function" is a bad subagent job.

tools is a restriction, not advice. An agent without `Edit` and `Write` **cannot** modify a file, whatever it decides. Same for the three isolation models, `fork` copies the parent window, `teammate` starts cold, `worktree` gets its own checkout.

Skills are routed by their description alone. No embedding, no classifier. The harness puts every skill's one-line description in the prefix, the model reads them as ordinary text, and emits `Skill(skill="...")`. **The description is the product**, a perfect body behind a vague line never loads.

A prompt is a probabilistic request. A hook is a guarantee. "Always run the formatter after editing" written in a prompt works most of the time. The same rule as a post-edit hook works every time and costs zero tokens. Judgement goes in the prompt, guarantees go in hooks.

Command intake

Not every line you type reaches the model. The harness reads the raw input first, and the split is binary.

```
> /clear          harness only, 0 tokens, the model is never called
> !git status     the shell runs it, only the output enters the window
> @src/mt/model.py the file is attached to the next turn
> /code-review high SKILL.md expands into a prompt, the loop starts
> why is the loss noisy plain text, the loop starts
```

`!git status` is **not** a `Bash` tool call. The model never decided it, never saw it, and no permission prompt fires. You ran it, the harness just shows the output to both of you.

How it recognises and runs them

Step	What happens	Example
1. First character	lexical dispatch, no model involved	<code>/</code> , <code>!</code> , <code>@</code> , or text
2. Start of line only	position 0 or it is prose	<code>/clear</code> fires, <code>run /clear later</code> does not
3. Lookup	built-ins first, then project skills, then user skills	a project <code>/deploy</code> shadows a personal one
4a. Built-in	the harness acts, the turn ends, nothing is appended	<code>/context</code> prints and stops, 0 tokens
4b. Skill	the markdown body becomes the prompt, arguments passed through	<code>/code-review high</code> → review instructions + <code>high</code>
5. Anything else	appended as a normal user turn	<code>why is the loss noisy</code>

 **Commands chain**, up to six skills on one line sharing the same arguments, because expansion is textual and concatenating two prompts is trivial.

 **They queue behind a running turn**, since injecting into a half-built window would corrupt it. Only `/status`, `/tasks` and `/usage` jump the queue, they read state without touching the conversation.

 **Custom commands and skills are the same thing now.** `.claude/commands/deploy.md` and `.claude/skills/deploy/SKILL.md` both produce `/deploy`.

Making the model recognise `/clear` would be slower (a full forward pass for five characters), more expensive, and probabilistic. **It is the hook rule one layer earlier.**

▼ The full command list (v2.1.220, click to expand)

Context and history

Command	Utility
<code>/clear [name]</code>	wipe the window, start a fresh conversation
<code>/compact [instructions]</code>	summarise the tail now instead of waiting for the budget
<code>/context [all]</code>	show what is really occupying the window right now
<code>/autocompact [auto <tokens>]</code>	set the budget that triggers compaction
<code>/rewind [count]</code>	roll code and conversation back to a checkpoint
<code>/recap</code> , <code>/view-session</code>	one-line summary, or the raw session context
<code>/export [file]</code> , <code>/copy [N]</code>	dump the conversation, or copy the Nth answer

Session shape

Command	Utility
<code>/resume [name]</code>	reopen an earlier conversation
<code>/branch [name]</code>	branch the current conversation in place
<code>/fork [prompt]</code>	copy the conversation into a new background session
<code>/background [prompt]</code>	detach this session as a background agent
<code>/subtask [prompt]</code>	hand one side task to a subagent
<code>/goal [condition clear]</code>	keep working until a condition holds
<code>/btw [question]</code>	ask something without polluting the history
<code>/tasks</code> , <code>/list-agents</code>	what runs in the background, who is reachable
<code>/teleport</code> , <code>/desktop</code> , <code>/mobile</code>	move the session to another surface

Model and cost

Command	Utility
<code>/model [name]</code>	switch model and save it as the default
<code>/fast [on off]</code>	faster output on the same Opus model, not a smaller one
<code>/usage (/cost)</code>	tokens burnt and money spent
<code>/advisor [model off]</code>	consult a second model on the side
<code>/status</code>	session status

Capability wiring

Command	Utility
<code>/agents</code>	manage subagent definitions
<code>/hooks</code>	inspect which hooks fire on which tool event
<code>/mcp [reconnect enable disable]</code>	manage MCP servers and their OAuth
<code>/plugin , /reload-plugins , /unload-plugin</code>	manage plugins
<code>/memory</code>	edit <code>CLAUDE.md</code> and the auto memory
<code>/permissions (/allowed-tools)</code>	allow, ask and deny rules
<code>/trust [path all none]</code>	which workspaces are trusted

Workspace and files

Command	Utility
<code>/add-dir <path></code>	give the session another readable directory
<code>/cd <path></code>	move the working directory
<code>/diff , /ide</code>	diff viewer, IDE integration
<code>/artifacts</code>	list, attach or open published artifacts

Plumbing

Command	Utility
<code>/help</code>	the list itself, and the only source of truth that never goes stale
<code>/config</code> (<code>/settings</code>)	settings UI, or <code>key=value</code> directly
<code>/login</code> , <code>/logout</code> , <code>/privacy-settings</code>	account and privacy
<code>/upgrade</code> , <code>/heapdump</code>	update, or dump the heap when something is wrong
<code>/theme</code> , <code>/color</code> , <code>/focus</code> , <code>/keybindings</code>	appearance and input
<code>/exit</code> (<code>/quit</code>)	leave

Skill commands (these do call the model)

Command	Utility
<code>/code-review</code> [level] [--fix] (<code>/review</code>)	review the diff or a PR for bugs and cleanups
<code>/security-review</code> [path PR]	check the diff for vulnerabilities
<code>/plan</code> [description]	enter plan mode and structure a complex task
<code>/init</code>	write the project's <code>CLAUDE.md</code>
<code>/debug</code> [description]	debug logging and runtime diagnosis
<code>/deep-research</code> <question>	fan out web research, synthesise a cited report
<code>/batch</code> <instruction>	large-scale parallel changes through worktrees
<code>/loop</code> [interval] [prompt]	run a prompt on a schedule
<code>/dataviz</code> , <code>/design-sync</code>	charts and dashboards, design system sync
<code>/claude-api</code> , <code>/doctor</code> , <code>/insights</code>	API references, install diagnosis, usage report
<code>/fewer-permission-prompts</code>	build a permission allowlist from your transcripts
your own	every skill in <code>.claude/skills/</code> gets a <code>/name</code>

⚠️ Commands move between families as features graduate, so `/help` is the listing that is never wrong.

Why it works, in five lines

- 1 The model decides, the harness executes.** Failures stay attributable.
- 2 Failure is information.** Errors go back into the window, and recovery becomes reasoning.

- 3 **The window is the product.** Append-only, because the cache says so.
- 4 **Pay only for what is used.** Catalogue in context, content on demand.
- 5 **Guarantees are code.** Anything that must always happen is a hook, never a sentence.

None of these need a frontier model. They need a loop that respects the window, which is why this shape is worth building around a small local model too.

Repository structure

```
├── assets/
│   ├── banner.svg
│   └── architecture.svg      # the map, one column per mechanism
├── scripts/
│   └── md_to_pdf.py          # renders this README to PDF via headless Edge
├── src/harness/              # 🚧 planned, not written yet
│   ├── loop.py               # the fifteen lines, and nothing else
│   ├── providers/            # base, local_mt (drives modern-transformer), anthropic
│   ├── context/              # window assembly, compaction, disclosure
│   ├── tools/                # registry, parallel executor, filesystem tools
│   ├── agents.py             # fork, teammate, worktree
│   └── hooks.py              # pre and post tool interception
├── skills/                   # 🚧 planned, one folder per skill
├── LICENSE
└── README.md
```

Run it on Your PC

🚧 **There is nothing to run yet.** The README is the current state, the package is the plan. Cloning still helps, the diagram reads better locally than on GitHub (the `Inter` import is blocked by GitHub's sandbox).

```
git clone https://github.com/Thibault-GAREL/LLM_harness.git
cd LLM_harness

python scripts/md_to_pdf.py      # regenerate README.pdf after any edit
```

Once `src/harness/` exists, the intended setup is below.

```
python -m venv .venv # if you don't have a virtual environment
source .venv/bin/activate # Linux / macOS
.venv\Scripts\activate   # Windows

pip install torch pydantic pyyaml rich

python -m harness --provider local_mt --checkpoint path/to/mt_weights.pt
```

⚠️ The `local_mt` provider needs a **CUDA-compatible GPU** to be usable interactively. On 6 GB of VRAM a small `modern-transformer` checkpoint fits, a large one does not, and the `anthropic` provider keeps the loop identical either way.

📖 Inspiration / Sources

This project is a study, so the sources matter more than usual:

- 📄 [The Claude Code source leak, 512 000 lines and a missing .npmignore](#) (Layer5, the reference post-mortem)
- 🔍 [Diving into Claude Code's source code](#) (Engineer's Codex, the architecture breakdown)
- 🦋 [claw-code](#) (the clean-room Python and Rust rewrite published after the leak)
- 📄 [Building Effective AI Agents](#) (Anthropic engineering)
- 📄 [Effective context engineering for AI agents](#) (Anthropic engineering)
- 🧬 [modern-transformer](#) (my own model, the provider this harness is built around)
- 🤖 [Language Models from Scratch](#) (where the series starts, a bigram model and a 2017 Transformer)

⚖️ No leaked proprietary source is reproduced here. Everything above describes **architecture and design principles**, which is exactly the part that transfers to a harness of your own.

Code created by me 🤖, Thibault GAREL - [Github](#)